



Learning curves from the archived runs

Each objective has its own scale. Lower training loss does not establish task success.

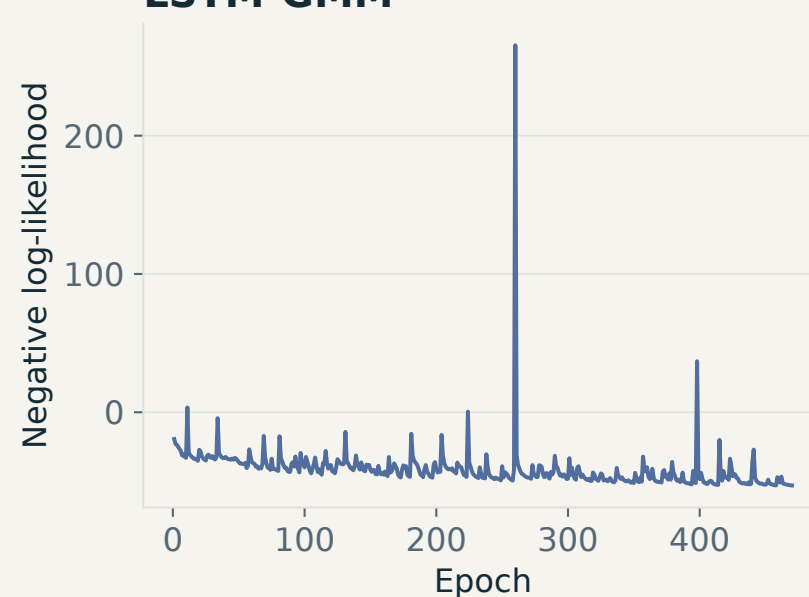
Diffusion Policy



IBC



LSTM-GMM



Last logged loss per epoch; repeated global steps resolved by keeping the last record. No smoothing.

DP / IBC: 450 epochs. Archived LSTM-GMM log: 470 epochs, with resume history; the paper reports 450.